

RAG Fundamentals

Retrieval-Augmented Generation (RAG) grounds LLM answers in external knowledge.

Ingestion steps: parse document, chunk text, embed chunks, index in vector DB.

Retrieval steps: embed query, search vectors, rerank, assemble context.

Generation step: inject context and ask model to answer with evidence.